

NAAN MUDHALVAN PROJECT

PHASE 2 : INNOVATION

CUSTOMER CHURN PREDICTION

Team Members :

ANANYA K A – 2021115013

ANULATHA S K - 2021115014

ARADHYA M - 2021115015

SAKTHIVEL K – 2021115088

BALAJI P - 2021115323

Project Overview:

A customer churn prediction project is a common and valuable application of data analytics, particularly in the fields of marketing and customer relationship management. Churn prediction involves identifying customers who are likely to stop doing business with a company in the near future. By identifying these customers early, a company can take proactive measures to retain them, which is often more cost-effective than acquiring new customers. Here's an overview of the key steps involved in a customer churn prediction project:

1. Understanding the Problem:

- Define the problem: Clearly state the objective of your project, which is to predict customer churn.
- Understand the business context: Familiarize yourself with the industry, the company, and its specific customer churn challenges. This understanding will guide your analysis.

2. Data Collection:

- Gather data: Collect relevant data, such as customer information, transaction history, customer interactions, and any other data that might be indicative of churn.

3. Data Preprocessing:

- Data cleaning: Address missing values and outliers in the dataset.
- Feature engineering: Create or modify features that may be relevant for predicting churn. For example, you might calculate metrics like customer lifetime value or churn propensity.
- Data transformation: Encode categorical variables, scale numerical features, and perform other preprocessing steps.

4. Exploratory Data Analysis (EDA):

- Conduct EDA to gain insights into the data. Explore the distribution of variables and relationships between features and the target variable (churn).

5. Data Splitting:

- Divide your dataset into training and testing sets to evaluate the model's performance accurately.

6. Model Selection:

- Choose appropriate machine learning or predictive modeling algorithms. Common choices include logistic regression, decision trees, random forests, support vector machines, and gradient boosting.

7. Model Training:

- Train your selected model(s) on the training data. Consider using techniques like cross-validation for hyperparameter tuning and model selection.

8. Model Evaluation:

- Assess the model's performance using appropriate evaluation metrics, such as accuracy, precision, recall, F1-score, and AUC-ROC. These metrics will help you understand how well your model predicts customer churn.

9. Model Interpretation:

- Interpret the model to understand which features are most influential in predicting churn. This insight can guide business decisions.

10. Deployment:

- If the model performs well, deploy it into a production environment, such as a customer relationship management (CRM) system, to make real-time predictions.

11. Monitoring and Iteration:

- Continuously monitor the model's performance in a production environment and update it as needed. Customer behavior and patterns may change over time.

12. Report and Communication:

- Summarize your findings, insights, and model performance in a clear and concise report. Present your results to relevant stakeholders in the organization.

13. Actionable Insights:

- Provide actionable recommendations based on your analysis. This could include strategies for retaining at-risk customers, personalized marketing, or other initiatives to reduce churn.

Model Selection :

Model Selection is a crucial step in the process of building a customer churn prediction model. It involves choosing an appropriate algorithm or method to train a model that can effectively predict customer churn. Below, I'll discuss some commonly used models for churn prediction and justify their selection.

1. Logistic Regression:

- Justification: Logistic regression is a simple and interpretable model that is often a good starting point for churn prediction. It's a binary classification algorithm that models the probability of a customer churning. Here's why it's a good choice:

- Interpretability: Logistic regression provides clear insights into the impact of each feature on the probability of churn. This interpretability is important for understanding the drivers of churn.

- Simplicity: Logistic regression is relatively simple and computationally efficient, making it a good choice for quick prototyping and model building.

- Works well with small to moderate-sized datasets: If you have a limited amount of data, logistic regression can still perform well without overfitting.

2. Decision Trees:

- Justification: Decision trees are another common choice for churn prediction. They are a non-linear model that can capture complex interactions between features.

- Interpretability: Although not as interpretable as logistic regression, decision trees can still provide insights into the features that influence churn. You can easily visualize the tree structure.

- Handles non-linear relationships: Decision trees can handle non-linear relationships between features and the target variable, which may be important in churn prediction where the factors affecting customer decisions can be complex.

3. Random Forests:

- Justification: Random forests are an ensemble learning technique that combines multiple decision trees, providing improved predictive performance and robustness.

- Improved predictive performance: Random forests reduce overfitting compared to individual decision trees and generally provide better generalization to unseen data.

- Feature importance: Random forests can measure feature importance, helping you understand which features are most relevant in predicting churn.

- Robustness: They are less sensitive to outliers and noise in the data, which can be beneficial in real-world scenarios.

4. Gradient Boosting (e.g., XGBoost or LightGBM):

- Justification: Gradient boosting algorithms like XGBoost or LightGBM have become increasingly popular for churn prediction due to their high predictive accuracy and versatility.
- State-of-the-art performance: These algorithms often achieve top performance in various machine learning competitions and churn prediction challenges.
- Handling imbalanced data: Churn prediction datasets are often imbalanced, with more non-churn cases. Gradient boosting can handle imbalanced data effectively by weighting classes or using sampling techniques.
- Automatic feature selection: Gradient boosting models can automatically select important features, reducing the need for extensive feature engineering.

5. Support Vector Machines (SVM):

- Justification: SVMs are a good choice when you want to find a hyperplane that maximizes the margin between churn and non-churn cases in a high-dimensional space.
- Effective in high-dimensional spaces: Churn prediction often involves many features, and SVMs are effective in such scenarios.
- Good generalization: SVMs aim to find the best separating hyperplane, which often results in good generalization to unseen data.

- Flexibility in kernel selection: SVMs can use various kernel functions to capture complex relationships between features.

The choice of model should be driven by the specific characteristics of your dataset, our computing resources, and your objectives. It's often a good practice to start with simpler models like logistic regression and then progressively explore more complex models to see if they offer improved performance. Additionally, one can use techniques like cross-validation to compare model performance and select the one that works best for our churn prediction problem.

XGBoost (Extreme Gradient Boosting):

It is a popular and powerful choice for a customer churn prediction project for several compelling reasons. It has gained prominence in the machine learning community and is often a top choice for various predictive modeling tasks, including churn prediction. Here's why XGBoost is a better choice for this project:

1. High Predictive Performance:

- XGBoost is known for its outstanding predictive performance. It consistently ranks among the top-performing algorithms in machine learning competitions and real-world applications. In the context of customer churn prediction, its

ability to capture complex, non-linear relationships between features can lead to more accurate predictions.

2. Regularization Techniques:

- XGBoost incorporates L1 (Lasso) and L2 (Ridge) regularization techniques, which help prevent overfitting. This is particularly useful when dealing with noisy or high-dimensional datasets common in churn prediction.

3. Handles Imbalanced Data:

- Customer churn datasets often suffer from class imbalance, where the number of customers who don't churn far outweighs those who do. XGBoost can handle imbalanced datasets effectively by assigning different weights to classes or using sampling techniques.

4. Feature Importance:

- XGBoost provides a feature importance score, allowing you to identify the most influential features in predicting churn. This is essential for understanding the drivers of churn and can inform business decisions.

5. Efficient Parallel and Distributed Computing:

- XGBoost is optimized for efficiency and can take advantage of parallel and distributed computing. This is valuable when

working with large datasets, as it can significantly reduce training time.

6. Flexibility in Loss Functions:

- XGBoost supports a variety of loss functions, making it adaptable to different types of classification problems. You can select a loss function that aligns with the specific characteristics of your churn prediction problem.

7. Gradient Boosting:

- XGBoost is a gradient boosting algorithm, which means it builds a series of decision trees sequentially, learning from the errors of previous trees. This approach often leads to improved model accuracy, as the model adjusts its predictions to focus on the most challenging cases.

8. Easy to Tune:

- XGBoost has numerous hyperparameters that can be fine-tuned to optimize its performance for your specific dataset. Techniques like cross-validation can help find the best set of hyperparameters.

9. Growing Ecosystem:

- XGBoost has a strong and growing ecosystem, with extensive community support, documentation, and various

libraries and tools built around it. This makes it easier to integrate into your project and access resources for troubleshooting.

10. Deployment:

- XGBoost models can be deployed in various production environments, making it suitable for real-time churn predictions, which is often a critical requirement in customer retention strategies.

While XGBoost is a powerful choice, it's important to note that no single model is universally the best for every situation. It's advisable to experiment with multiple models, including XGBoost, and rigorously evaluate their performance on our specific dataset. The decision to use XGBoost should be based on how well it performs in our churn prediction task, considering the factors mentioned earlier, including data characteristics, computational resources, and our objectives.